

Deep Reinforcement Learning and Control



Multigoal RL

CMU 10-703

Katerina Fragkiadaki



So far we train one policy/value function per task, e.g., win the game of Tetris, win the game of Go, reach to a **particular** location, put the green cube inside the gray bucket, etc.



Universal value function Approximators

$$V(s; \theta) \rightarrow V(s, g; \theta)$$

$$\pi(s; \theta) \rightarrow \pi(s, g; \theta)$$

- All methods we have learnt so far can be used.
- At the beginning of an episode, **we sample not only a start state but also a goal g** , which stays constant throughout the episode
- The experience tuples should contain the goal.

$$(s, a, r, s') \rightarrow (s, g, a, r, s')$$

Universal value function Approximators

$$\begin{array}{ccc} V(s, \theta) & \rightarrow & V(s, g; \theta) \\ \pi(s; \theta) & \rightarrow & \pi(s, g; \theta) \end{array}$$

What should be my goal representation?

The goal representation is usually the same as your state representation.
Usually one of the two:

- **Manual/oracle:** 3d centroids of objects, robot joint angles and velocities, 3d location of the gripper, etc.
- **Learnt:** Some feature encoding over images directly

Hindsight Experience Replay

Marcin Andrychowicz^{*}, Filip Wolski, Alex Ray, Jonas Schneider, Rachel Fong, Peter Welinder, Bob McGrew, Josh Tobin, Pieter Abbeel[†], Wojciech Zaremba[†]
OpenAI

Main idea: use failed executions under one goal g , as successful executions under an alternative goal g' (which is where we ended at the end of the episode).



Hindsight Experience Replay

Marcin Andrychowicz*, Filip Wolski, Alex Ray, Jonas Schneider, Rachel Fong, Peter Welinder, Bob McGrew, Josh Tobin, Pieter Abbeel[†], Wojciech Zaremba[†]
OpenAI

Main idea: use failed executions under one goal g , as successful executions under an alternative goal g' (which is where we ended at the end of the episode).



Hindsight Experience Replay

Algorithm 1 Hindsight Experience Replay (HER)

Given:

- an off-policy RL algorithm \mathbb{A} ,
- a strategy \mathbb{S} for sampling goals for replay,
- a reward function $r : \mathcal{S} \times \mathcal{A} \times \mathcal{G} \rightarrow \mathbb{R}$.

▷ e.g. DQN, DDPG, NAF, SDQN

▷ e.g. $\mathbb{S}(s_0, \dots, s_T) = m(s_T)$

▷ e.g. $r(s, a, g) = -[f_g(s) = 0]$

▷ e.g. initialize neural networks

Initialize \mathbb{A}

Initialize replay buffer R

for episode = 1, M **do**

Sample a goal g and an initial state s_0 .

for $t = 0, T - 1$ **do**

Sample an action a_t using the behavioral policy from \mathbb{A} :

$$a_t \leftarrow \pi_b(s_t || g)$$

▷ $||$ denotes concatenation

Execute the action a_t and observe a new state s_{t+1}

end for

for $t = 0, T - 1$ **do**

$$r_t := r(s_t, a_t, g)$$

Store the transition $(s_t || g, a_t, r_t, s_{t+1} || g)$ in R

▷ standard experience replay

Sample a set of additional goals for replay $G := \mathbb{S}(\text{current episode})$

for $g' \in G$ **do**

$$r' := r(s_t, a_t, g')$$

Store the transition $(s_t || g', a_t, r', s_{t+1} || g')$ in R

← G : the states of the current episode

▷ HER

end for

end for

for $t = 1, N$ **do**

Sample a minibatch B from the replay buffer R

Perform one step of optimization using \mathbb{A} and minibatch B


end for

end for

← Usually as additional goal we pick the goal that this episode achieved, and the reward becomes non zero

Hindsight Experience Replay

--- DDPG — DDPG+count-based exploration — DDPG+HER — DDPG+HER (version from Sec. 4.5)



Visual Reinforcement Learning with Imagined Goals

Ashvin Nair*, Vitchyr Pong*, Murtaza Dalal, Shikhar Bahl, Steven Lin, Sergey Levine


University of California, Berkeley

{anair17,vitchyr,mdalal,shikharbahl,stevenlin598,svlevine}@berkeley.edu


Main ideas:

- Train a generative model of images in an unsupervised manner (they used a VAE). The trained model can generate images by sampling latent codes from a gaussian distribution.
- Use that latent code as the state and goal representation
- Use L2 distance over latent codes as the (inverse of) reward function.
- Sample goals from that generative model for goal relabelling (augmenting experience)
- Retrain the generative model as the policy changes and the agent visits different parts of the state space


Learning Generative models of images



- Why simple gaussian noise suffices to create complex outputs?
- The neural net will transform it to a complex distribution!




$$z \sim \mathcal{N}(\mathbf{0}, I)$$



$$f(z) = \frac{z}{10} + \frac{z}{\|z\|}$$

Training Networks with Stochastic Units



Each sample z should give me a realistic image X once it passes through the neural network

We want to learn a mapping from z to the **output image X** , usually we assume a Gaussian distribution to sample every pixel from:

$$P(X|z; \theta) = \mathcal{N}(X|f(z; \theta), \sigma^2 \cdot I)$$

Let's maximize data likelihood. This requires an intractable integral, too many z s..

$$\max_{\theta} . \quad P(X) = \int P(X|z; \theta)P(z)dz$$


What if we forget that it is intractable and approximate it with few samples?

$$\min_{\theta} . \quad \sum_j -\log P(X_j) = - \sum_j \sum_{z_i \sim \mathcal{N}(\mathbf{0}, I)} \log P(X_j|z; \theta) = - \sum_j \sum_{z_i \sim \mathcal{N}(\mathbf{0}, I)} \|f(z_i; \theta) - X_j\|^2$$

(Q: do we know how to take gradients here?)

This is a bad approximation, except if we have a very large number of z s. Only few z s would produce after training reasonable X . How will we find the z s that produce good X ?

Training Networks with Stochastic Units



Each sample z should give me a realistic image X once it passes through the neural network

We want to learn a mapping from z to the **output image X** , usually we assume a Gaussian distribution to sample every pixel from:

$$P(X|z; \theta) = \mathcal{N}(X|f(z; \theta), \sigma^2 \cdot I)$$

Let's maximize data likelihood. This requires an intractable integral, too many z s..

$$\max_{\theta} . \quad P(X) = \int P(X|z; \theta)P(z)dz$$

What if we forget that it is intractable and approximate it with few samples?

$$\min_{\theta} . \quad \sum_j -\log P(X_j) = - \sum_j \sum_{z_i \sim \mathcal{N}(\mathbf{0}, I)} \log P(X_j|z; \theta) = - \sum_j \sum_{z_i \sim \mathcal{N}(\mathbf{0}, I)} \|f(z_i; \theta) - X_j\|^2$$

(Q: do we know how to take gradients here?)

$$\min_{\theta} . \quad - \sum_j \min_{z_i \sim \mathcal{N}(\mathbf{0}, I)} \|f(z_i; \theta) - X_j\|^2 \quad \text{K-best loss}$$

Deep Variational Inference

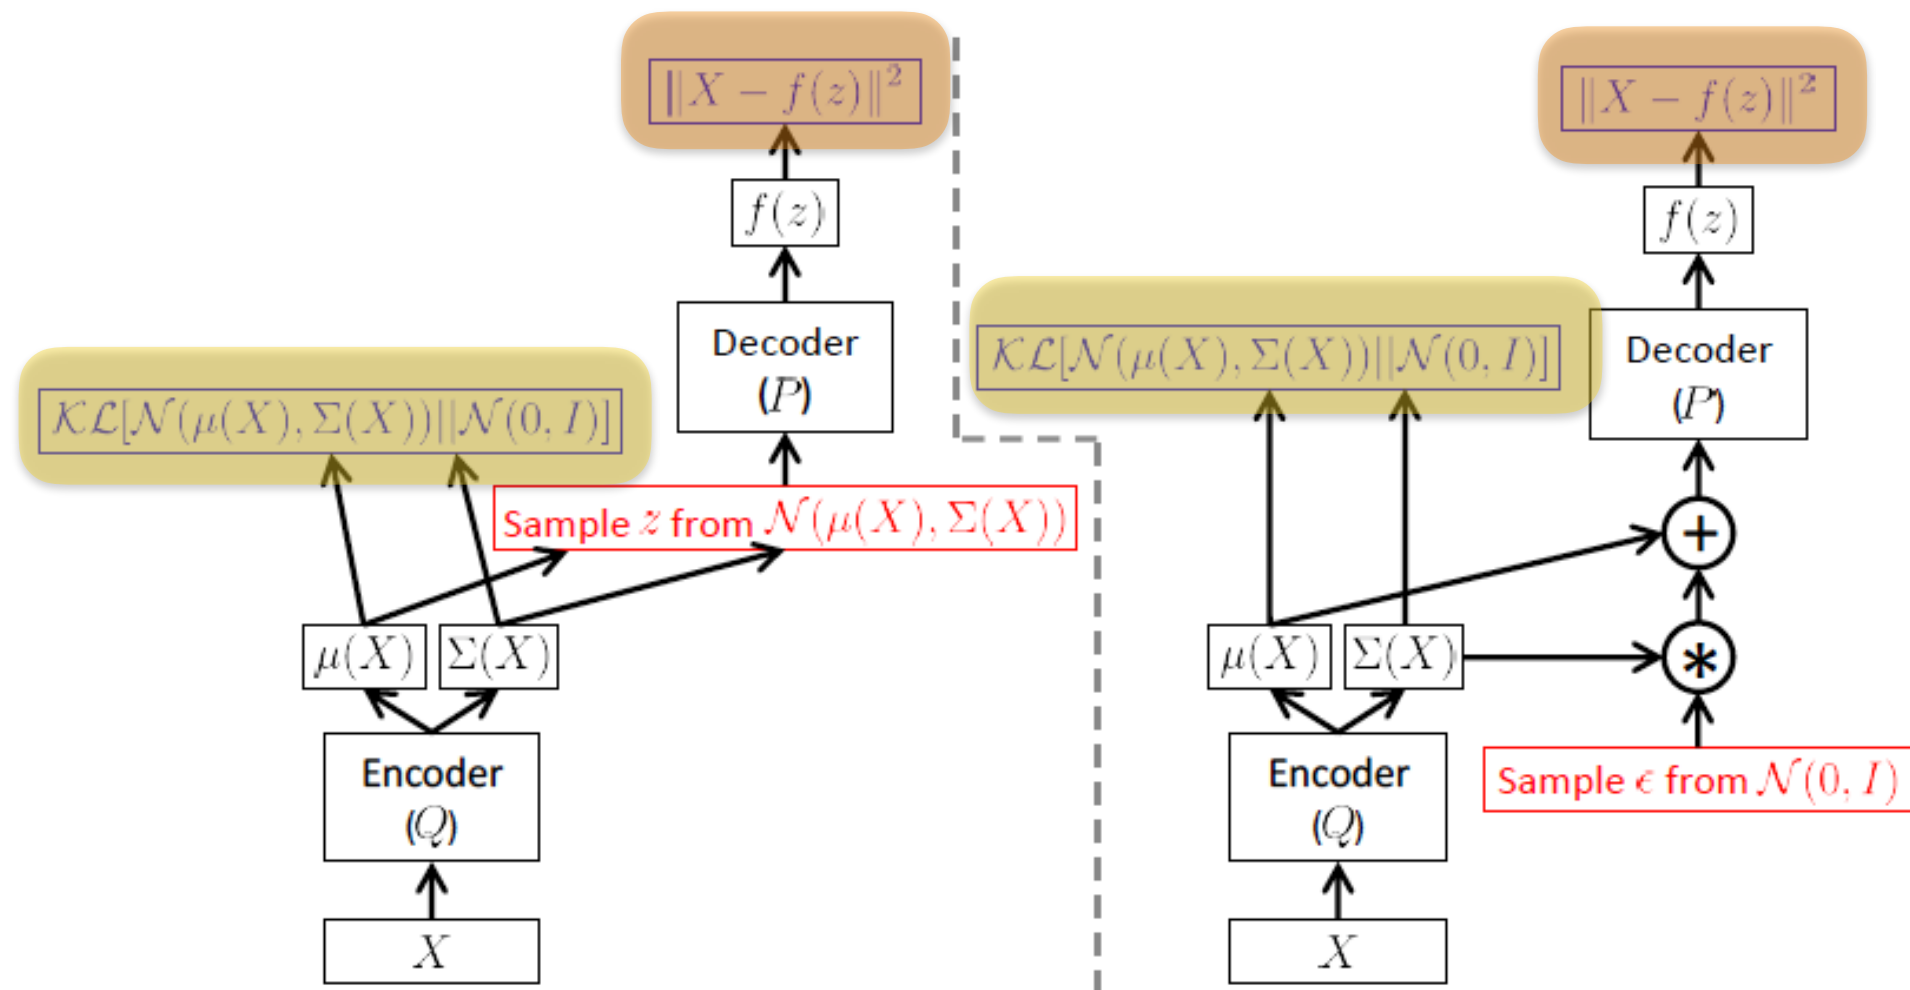
Let's consider sampling z s from an alternative distribution $Q(z)$ and try to minimize the KL between this (variational approximation) and the true posterior, $P(z|X)$. And because I can pick any distribution Q I like, I will also condition it on X to help inform the sampling.

$$\begin{aligned}D_{KL}(Q(z|X) || P(z|X)) &= \int Q(z|X) \log \frac{Q(z|X)}{P(z|X)} dz \\&= \mathbb{E}_Q \log Q(z|X) - \mathbb{E}_Q \log P(z|X) \\&= \mathbb{E}_Q \log Q(z|X) - \mathbb{E}_Q \log \frac{P(X|z)P(z)}{P(X)} \\&= \mathbb{E}_Q \log Q(z|X) - \mathbb{E}_Q \log \frac{P(X|z)P(z)}{P(X)} \\&= \mathbb{E}_Q \log Q(z|X) - \mathbb{E}_Q \log P(X|z) - \mathbb{E}_Q \log P(z) + \log P(X) \\&= D_{KL}(Q(z|X) || P(z)) - \mathbb{E}_Q \log P(X|z) + \log P(X)\end{aligned}$$

$$\min_{\phi, \theta} . \quad D_{KL}(\underbrace{Q(z|X; \phi)}_{\text{encoder}} || P(z)) - \mathbb{E}_Q \log \underbrace{P(X|z; \theta)}_{\text{decoder}}$$

Variational Autoencoder


From left to right: re-parametrization trick!




$$\min_{\phi, \theta} . \quad D_{KL}(\underbrace{Q(z | X; \phi)}_{\text{encoder}} || P(z)) - \mathbb{E}_Q \log \underbrace{P(X | z; \theta)}_{\text{decoder}}$$

Variational Autoencoder

At test time

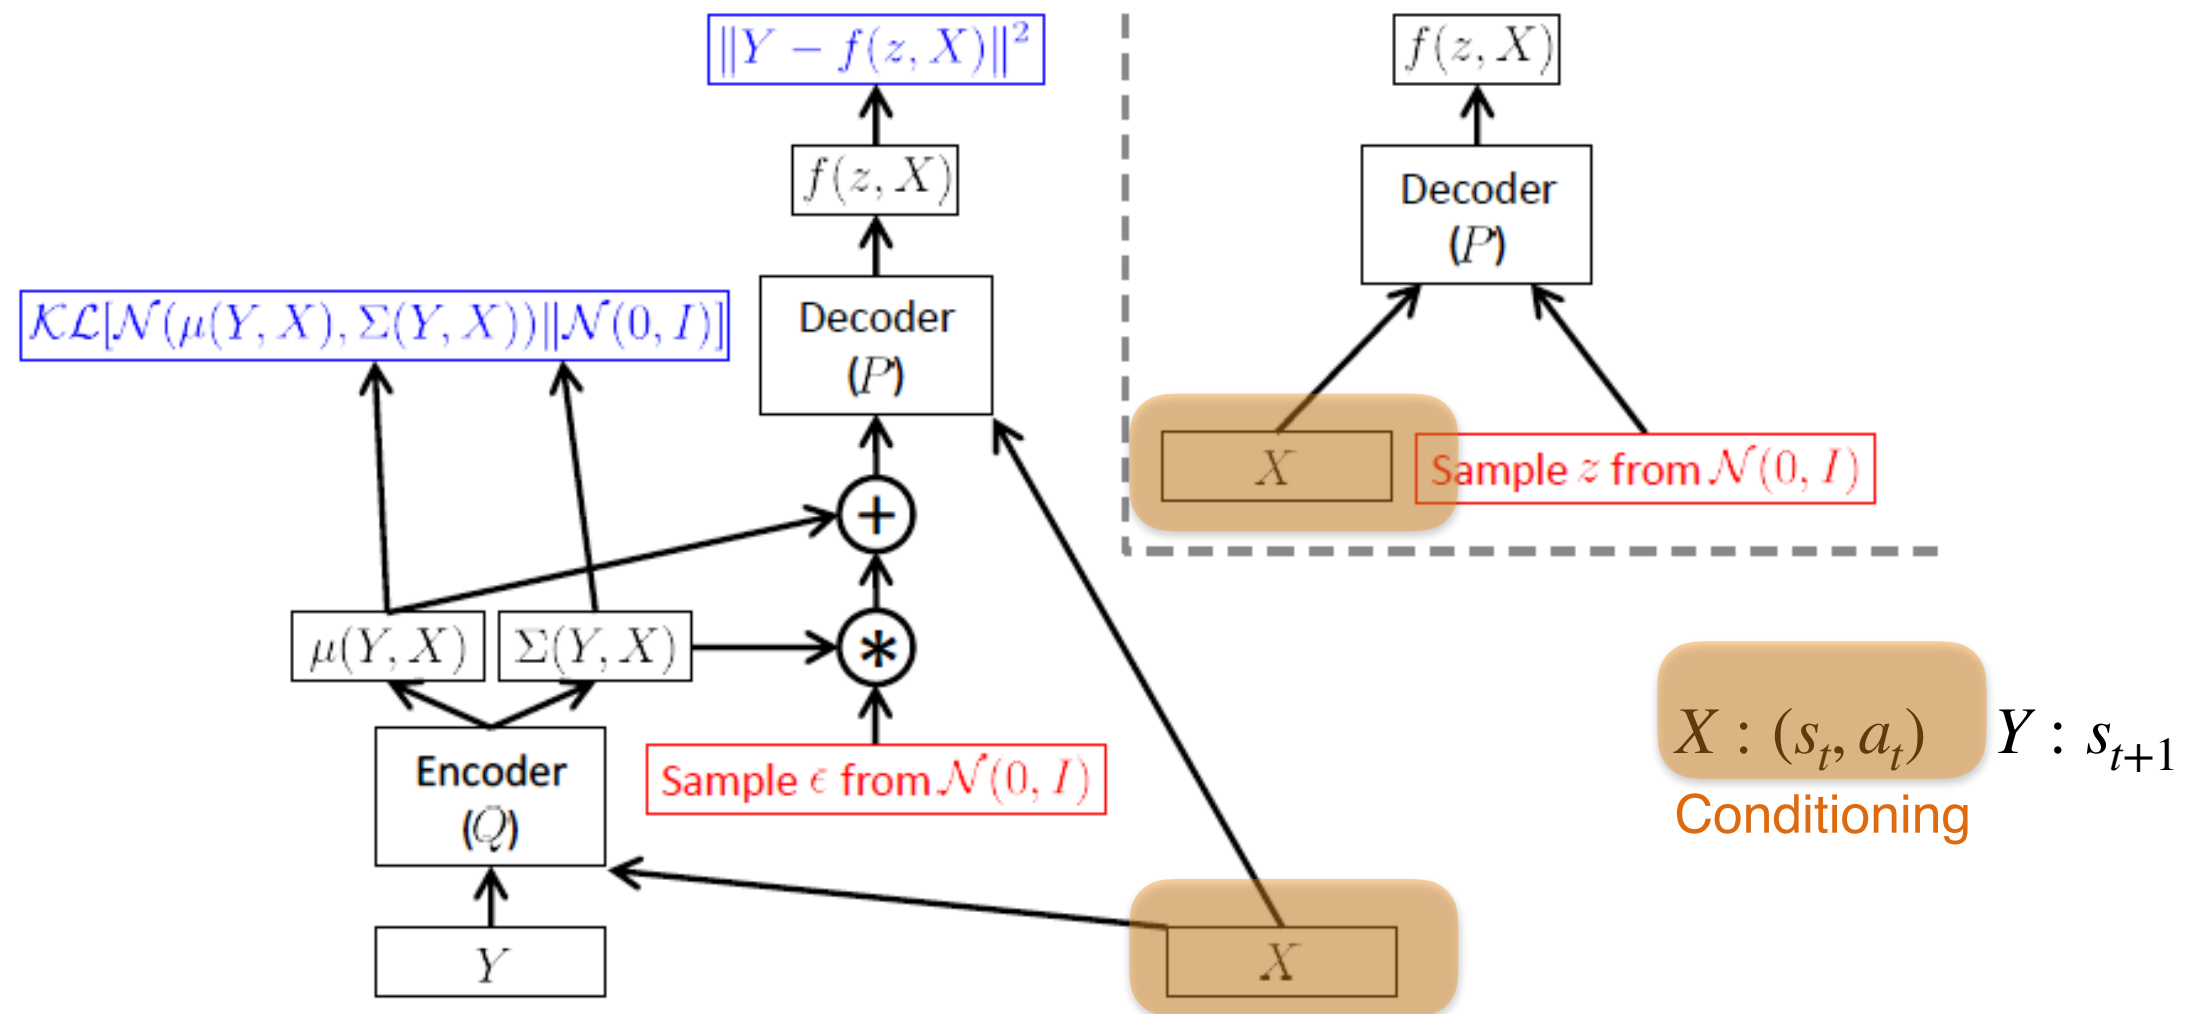


(a) Learned Frey Face manifold



(b) Learned MNIST manifold

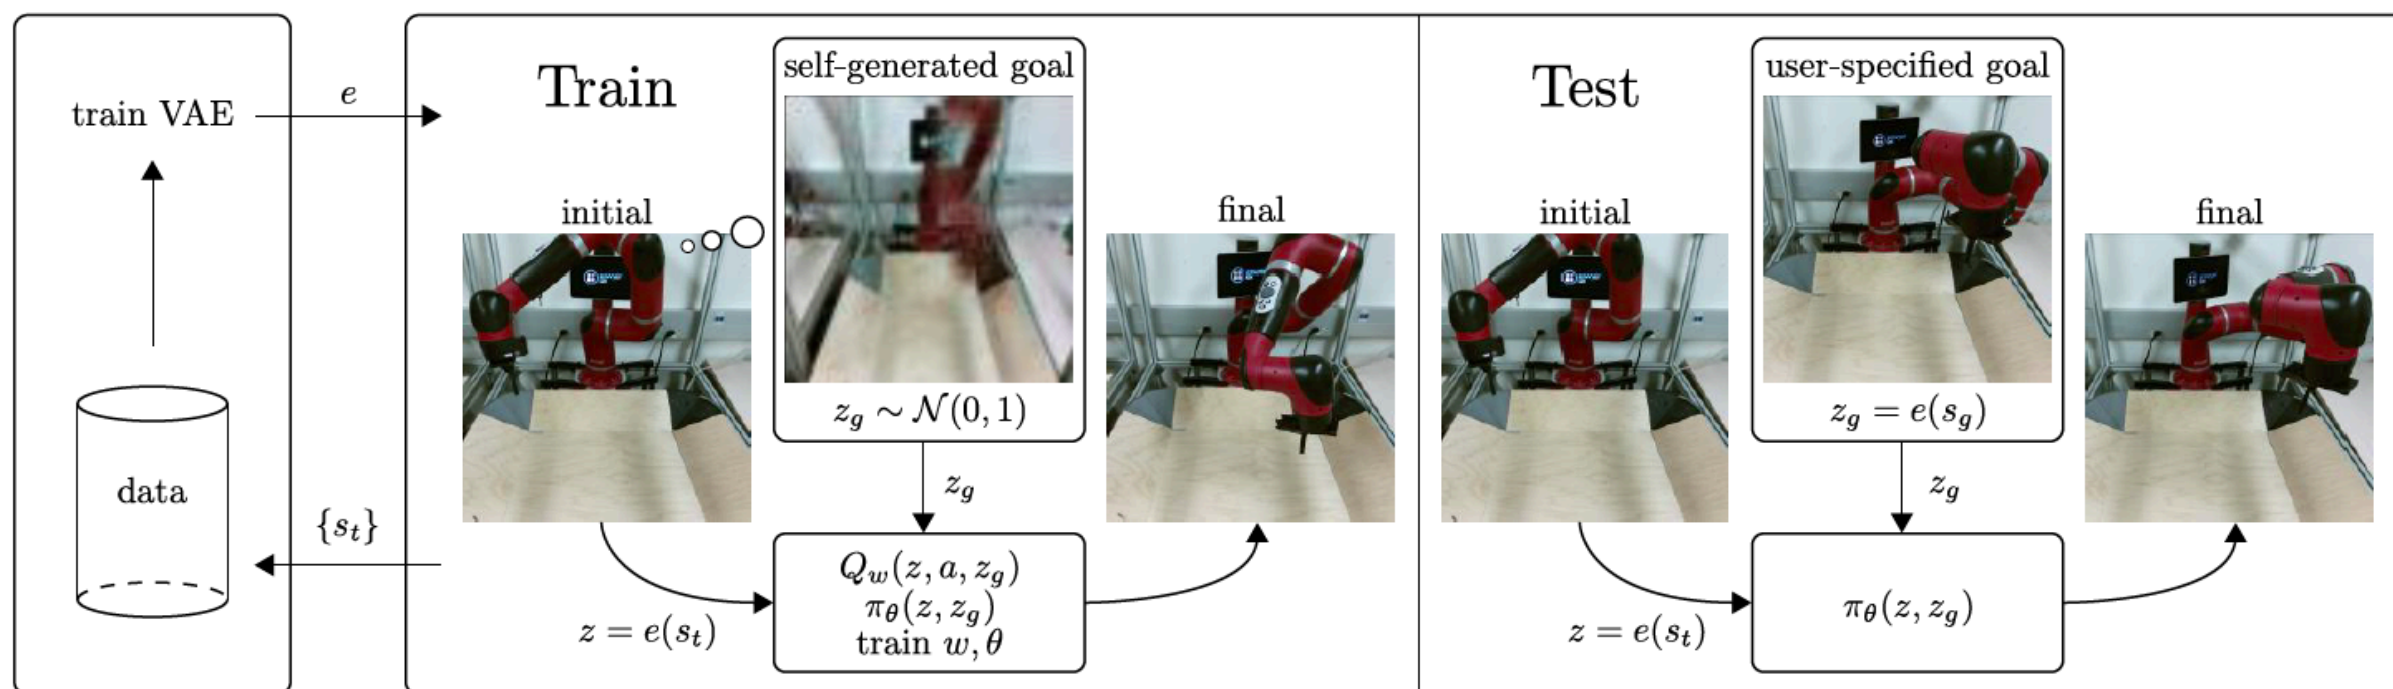
Conditional VAE



$$\min_{\phi} D_{KL}(Q(z | X, Y) || P(z | \mathcal{D})) = \min_{\phi} D_{KL}(Q(z | X, Y) || P(z)) - \mathbb{E}_Q \log P(\mathcal{D} | z)$$

Visual Reinforcement Learning with Imagined Goals

Ashvin Nair*, Vitchyr Pong*, Murtaza Dalal, Shikhar Bahl, Steven Lin, Sergey Levine
University of California, Berkeley
{anair17,vitchyr,mdalal,shikharbahl,stevenlin598,svlevine}@berkeley.edu



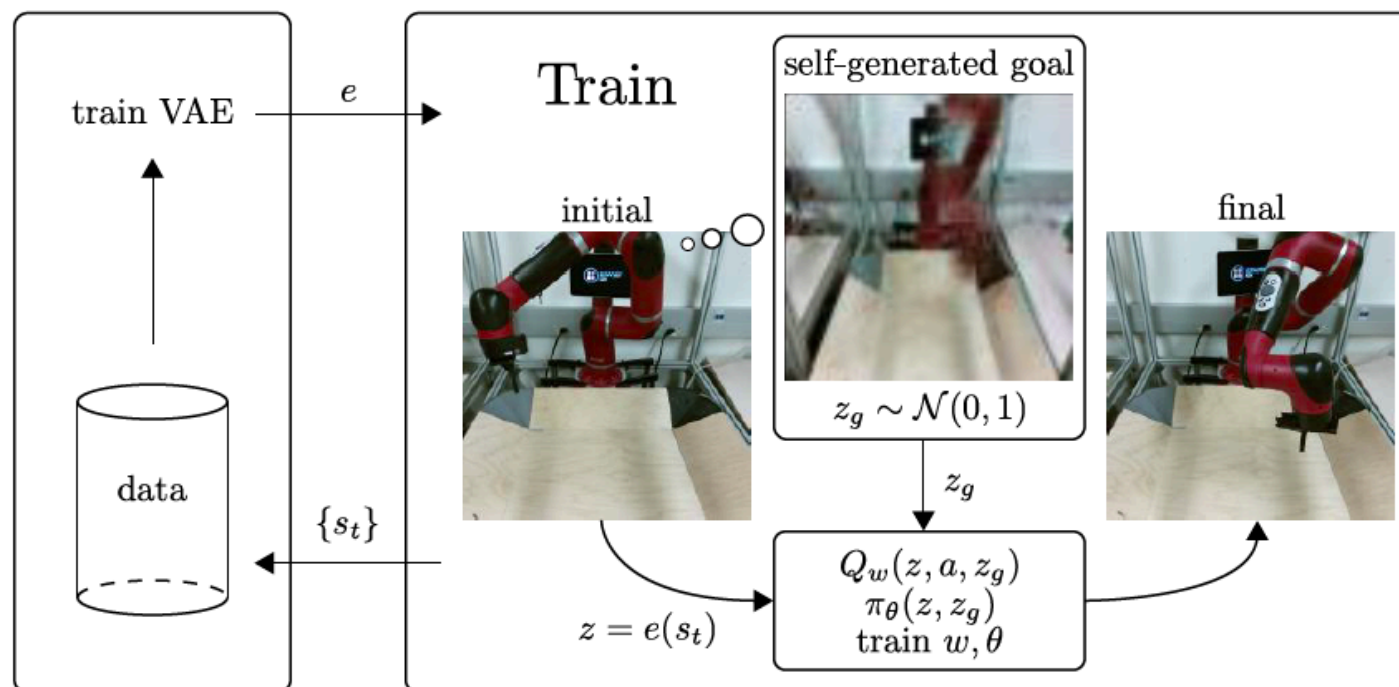
Visual Reinforcement Learning with Imagined Goals

Ashvin Nair*, Vitchyr Pong*, Murtaza Dalal, Shikhar Bahl, Steven Lin, Sergey Levine

University of California, Berkeley

{anair17,vitchyr,mdalal,shikharbahl,stevenlin598,svlevine}@berkeley.edu


At training time the agent imagines goals to reach by simply sampling codes (vectors) from the latent space.



Visual Reinforcement Learning with Imagined Goals

Ashvin Nair*, Vitchyr Pong*, Murtaza Dalal, Shikhar Bahl, Steven Lin, Sergey Levine
University of California, Berkeley
{anair17,vitchyr,mdalal,shikharbahl,stevenlin598,svlevine}@berkeley.edu

At test time, the human supplies a goal image which is encoded into a latent code by the trained encoder.



Algorithm 1 RIG: Reinforcement learning with imagined goals

Require: VAE encoder q_ϕ , VAE decoder p_ψ , policy π_θ , goal-conditioned value function Q_w .

1: Collect $\mathcal{D} = \{s^{(i)}\}$ using exploration policy.	9: Store (s_t, a_t, s_{t+1}, z_g) into replay buffer \mathcal{R} .
2: Train β -VAE on \mathcal{D} by optimizing (2).	10: Sample transition $(s, a, s', z_g) \sim \mathcal{R}$.
3: for $n = 0, \dots, N - 1$ episodes do	11: Encode $z' = e(s')$.
4: Sample latent goal from prior $z_g \sim p(z)$.	12: (Probability 0.5) replace z_g with $z'_g \sim p(z)$.
5: Sample initial state $s_0 \sim E$.	13: Compute new reward $r = - z' - z_g $.
6: for $t = 0, \dots, H - 1$ steps do	14: Minimize (1) using (z, a, z', z_g, r) .
7: Get action $a_t = \pi_\theta(e(s_t), z_g) + \text{noise}$.	15: end for
8: Get next state $s_{t+1} \sim p(\cdot s_t, a_t)$.	16: Fine-tune β -VAE every K episodes on mixture of \mathcal{D} and \mathcal{R} .
	17: end for

$$\mathcal{E}(w) = \frac{1}{2} ||Q_w(s, a, g) - (r + \gamma \max_{a'} Q_{\bar{w}}(s', a', g))||^2$$

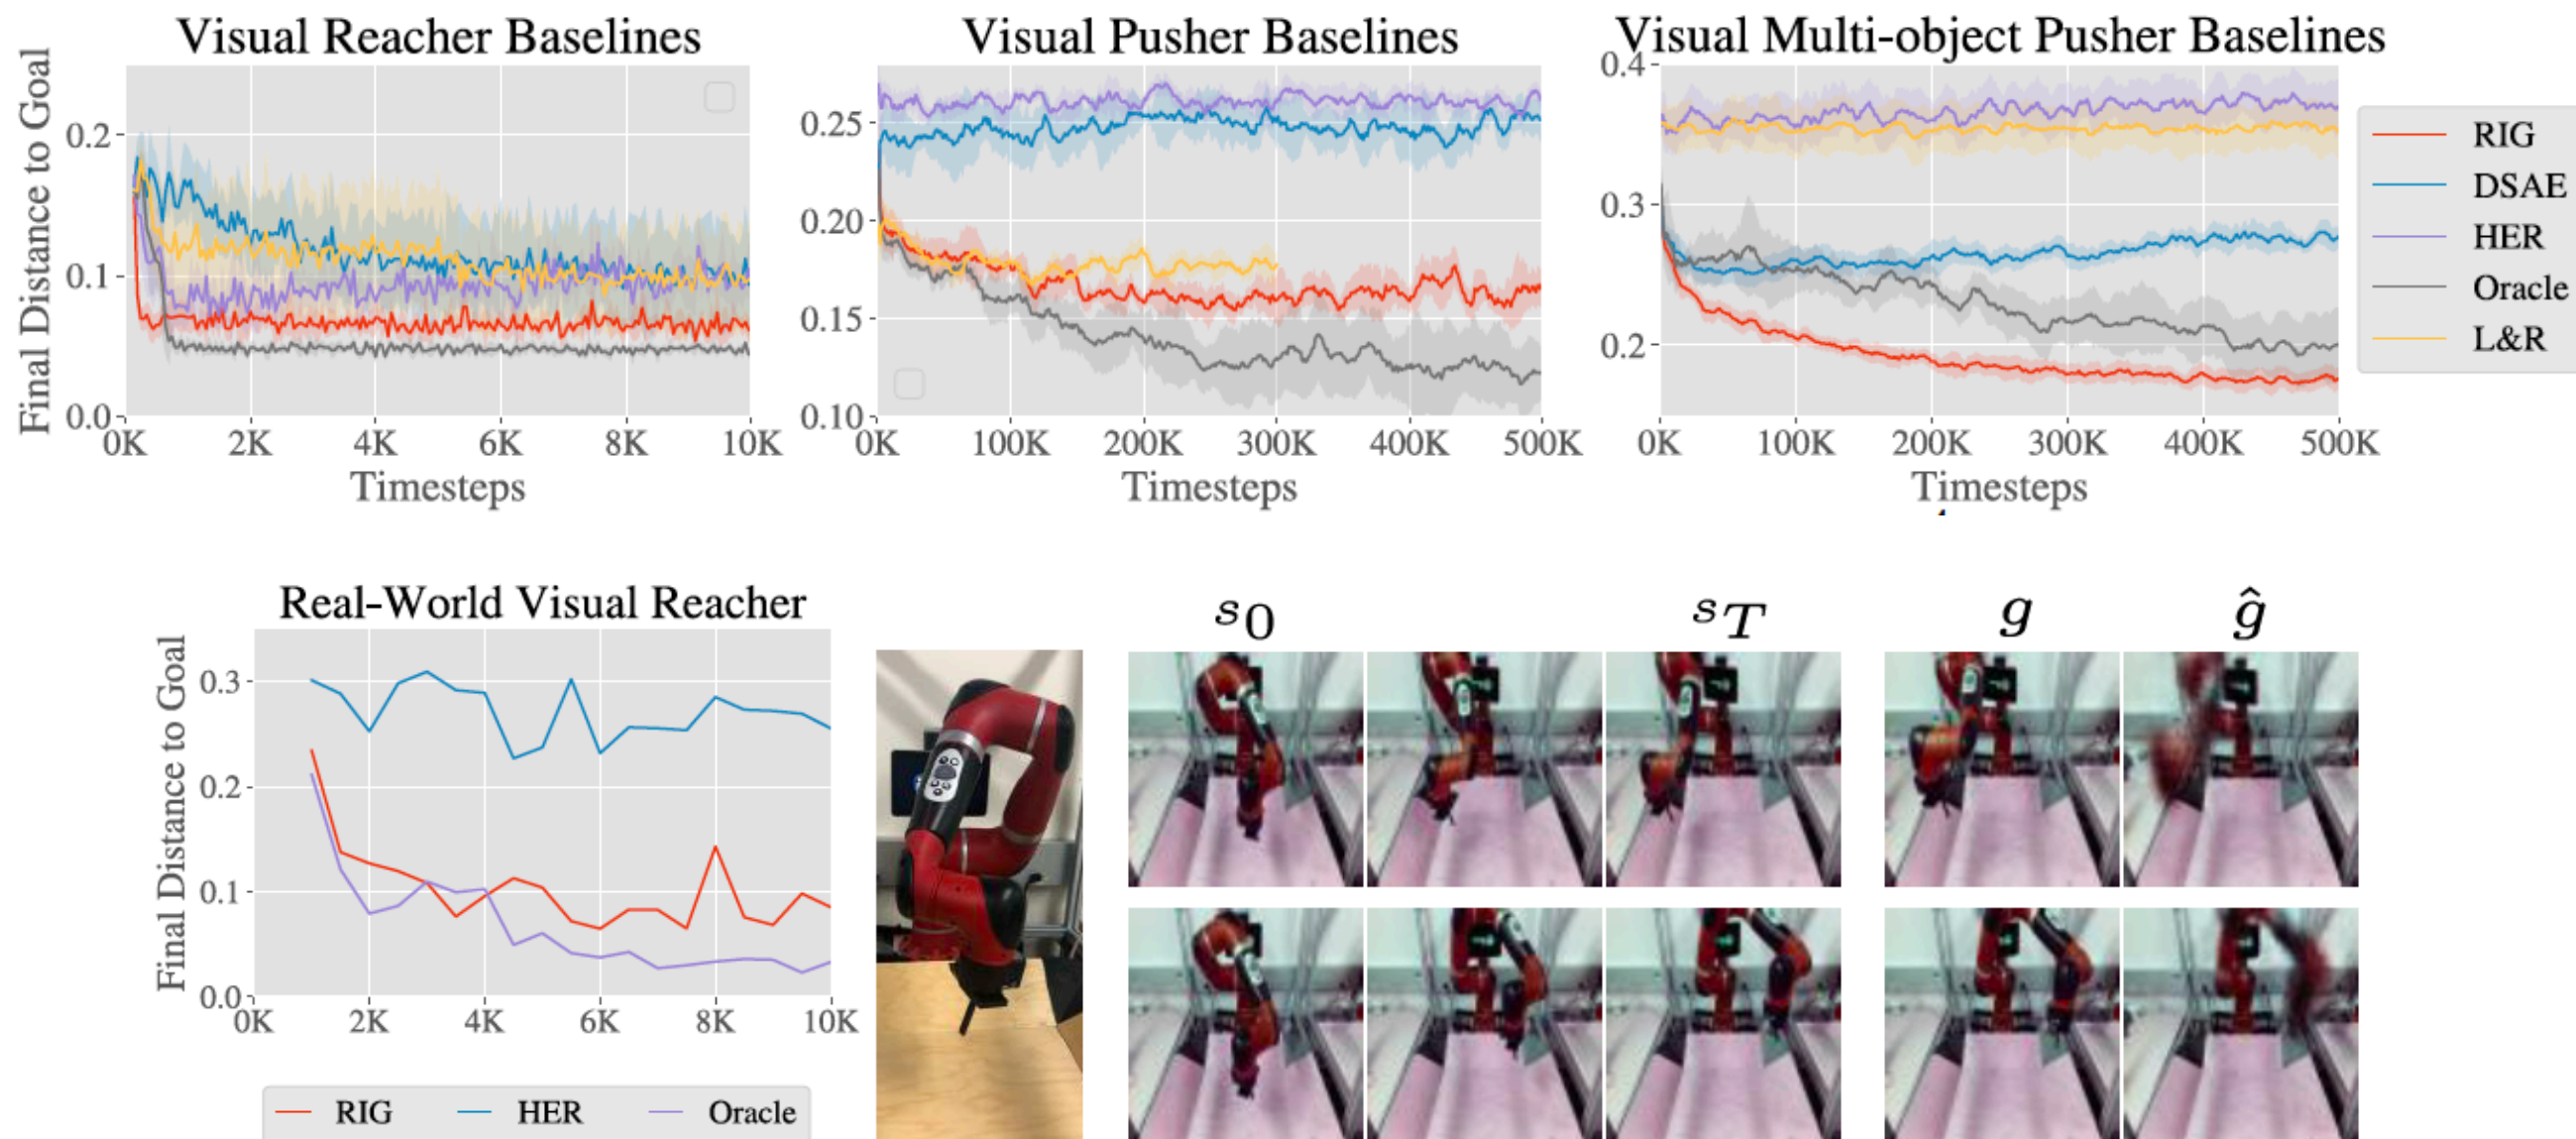


Figure 7: (Left) Our method compared to the HER baseline and oracle on a real-world visual reaching task. (Middle) Our robot setup is pictured. (Right) Test rollouts of our learned policy.

HER here is using L2 over images, that's a terrible (inverse of) reward function